

1. general information

- emphasis on market portfolios
- focus on statistics and simulations
- Genesis of risk
 - fat tails and skewness
 - non-linearities
 - wrong size
 - wrong horizon
 - Path dependency (your idea eventually works out but you are out of a job or without clients by then)
 - Model risk

2. statistics

- pdf, cdf, invcdf
- $p(y|\theta)$ as common in Bayesian statistics
- functions of random variables(rv) ★ see paper notes
- the next term detailed at "res"
- Stationarity:
 - Stationarity: the unconditional distribution of any observation can be defined independently of time
 - covariance stationary: the unconditional mean and variance exist
 - If we can isolate processes that are at least approximately stationary, we can build better models and decisions.
- The autocorrelation function $\rho(h)$: the correlation between y_t and y_{t+h}
- white noise if it is covariance stationary and uncorrelated
 - $e_t = u_t u_{t-1}^2$ is not strict white noise
 - see ps01
- NB! The unconditional distribution of y_t is theoretically defined only for stationary processes
- conditional distribution: $p(y_{t+h}|\Omega_t)$
 - eg ★ see paper notes
- skewness and kurtosis ★ see paper notes
- Some common measures of "risk":
 - Standard deviation
 - Semi-deviation,
 - mean absolute deviation (MD or MAD)
 - $\sum |y - \text{mean}(y)|$
 - semi-mad
 - VaR (value-at-risk)
 - ES (expected short-fall)
 - ES is a more relevant and better (better properties) measure of risk than VaR
 - ES is even harder to estimate accurately than VaR
 - heavy-tailed: ES is more relevant and also cannot be estimated accurately
 - Drawdown
 - Stress test (or "scenarios")
 - Convexity heuristics
 - ★ see paper notes

3. maximum likelihood

- OLS
 - however, normality is a very poor assumption for most financial returns
 - as the sample gets sufficiently large, MLE is better (less variance of errors)
- Maximum (log) likelihood
 - Probability density function

- Linear model with Gaussian errors
- ★ see paper notes
- OLS is asymptotically optimal in a well-defined sense for Gaussian
- Asymptotic distribution of ML estimators
 - $\theta_{ML} \sim N(\theta_0, I(\theta_0)^{-1})$ where $\theta_0 = \theta_{TRUE}$
 - $\hat{I}(\theta_{ML}) = - \sum_{i=1}^n \frac{\partial^2 I(\theta_{MLy_i})}{\partial \theta_{ML} \partial \theta_{ML}'}_{\theta_{ML}}$
- Maximum likelihood in practice
 - ...

4. normal, log-normal, CLT, fat tails

- z is risk-factor, and risk is on portfolio value $v = f(z)$
 - risk increase when:
 1. z is thick-tailed
 2. $f(\cdot)$ is non-linear (concave for right and convex for left) ★ see paper notes
 - an example of left-tail and thick-tail
 - z is normal
 - $y = \beta_0 + \beta_1 z + \beta_2 z^2 * I(z < 0)$
 - ◦ an example of thick-tailed distribution
 - $r_t = u_t u_{t-1}^2$ with normal u
- Why is the Gaussian distribution "normal"?
 - the sum of independent normals is still normal
 - central-limit-theorems
 - sums of sufficiently many non-normal r.v. are approximately normal.
 - the speed to normal depends on the individual random variables differ from normal
 - break or slow down CLT
 - infinite variance
 - not iid rv, but strong dependence
 - geometric average, CLT not hold (such as price)
 - but may be log normal
 - but can use log: ($\log(\text{product}) = \text{sum}(\log)$) \sim normal
- log normal
 - $mean = \exp(\mu + \frac{1}{2}\sigma^2)$
 - for small SD, approx to normal
 - right skewed and leptokurtic
 - Thinner tail if negatively affected by size
 - Thicker tail if positively affected by size
- extreme value theory

5. financial time series

- tails are much thicker(fat!)
 - in QQ plot, left lower and right higher
 - * skew < 0, kurtosis >> 3
- not continuous
 - jump
- **clustered volatility**: Strong and long-lived serial correlation of $|r|$ and r^2
- (many but not all) return series are strongly left-skewed
 - indexes are more skew than individual
 - monthly are more skew than daily
- CLT in most cases works very slowly
 - CLT: be normal from day->week->month->quarter->year
 - but slowly
 - **but: Skew in particular can initially grow with the horizon**

- reversion in some important variables (profit margins, valuations, oil prices...)
- often: use $\sigma(h) = \sqrt{h}\sigma(1)$ to compute $\text{VaR}(h)$, but:
 1. not normal
 2. more skewed at longer horizons
 - result underestimate

6. student-t

- portfolio with J assets
 - multivariate model
 - $R_t \sim MN(\mu, \Sigma)$
 - $R_{p,t} \sim N(w^T \mu, w^T \Sigma w)$
 - $R_{p,t} = \sum_{j=1}^J w_j R_{j,t}$
- $\chi^2(df = v) = z_1^2 + z_2^2 + \dots + z_n^2$
- $F(m, n) = \frac{\chi^2(n)/n}{\chi^2(m)/m}$
- student-t
 - $t(m, s, v) = N(m, s) * \sqrt{\frac{v}{\chi^2}}$
 - !!!! s^2 is not variance, but scale
 - Thicker tails for lower v
 - $\text{var}(t) = s^2 \frac{v}{v-2}$
 - pdf of t
- maximum Log-likelihood with student t errors
 - log-likelihood is extremely flat in v for $v > 30$ (hard to know the maximum one)
 - Using $\log(v)$ in estimation is an improvement
 - make sure to initialize the optimization from a low value of v, say, 4-8
 - cannot model skewness
- The t distribution can generate very fat tails but we must be aware of its limitations:
 - it cannot model skewness
 - some prefer to think of extreme events as generated by a different process
 - One way to move past both limitations is by using mixtures

7. optimization; mixtures of normals and t

- optimization (to maximum likelihood)
 - General advice
 - parameters have widely different size: standardize covariates and take logs of parameters
 - to transform parameters so that an unconstrained optimization: $\ln(\pi/(1-\pi))$
 - if have constraints: use Bayesian or quasi-Bayesian solution
 - deterministic
 - Newton-Raphson
 - NR has maximum efficiency in the best-case scenario (which we expect in large samples), but is fragile in less rosy scenarios
 - quasi-NR (BFGS) !!!
 - method = "BFGS" provides a much better default for most maximum likelihood problems
 - conjugate gradient (CG)
 - Nelder-Mead
 - stochastic (derivative-free, stochastic optimization)
 - gradient based stochastic approximation
 - random search
 - simulated annealing
 - can escape local modes
 - MUCH slower

- Analytical vs numerical derivatives
 - evaluating the log-likelihood numerically for small changes in parameters
 - p or 2p likelihood evaluations
 - Analytical derivatives will make estimation faster and more reliable
 - and the Hessian (and VCV (θ)) much more reliable.
 - but, need time to code, and some don't need hessian
- mixtures of normals and t
- Extensions
 - Regression model where the errors are a mixture of normals instead of normals
 - Each component has different regression parameters
 - π instead of π , with π a function of covariates x_t and/or an exogenous process.
- fewer and more complex is usually superior (when mixtures)
- Bonds returns are roughly symmetric
- Other returns are asymmetric, common structure
- Computing risk measures by simulation
 - a) mixture of t does a good job at reproducing the empirical VaR,
 - b) Gaussian VaR is OK at 0.95 but increasingly poor further out in the tail,
 - c) ES \gg VaR even at 0.999 due to low dof.

8. Modeling conditional volatility

- Volatility clustering: $|r|$ and r^2 autocorrelated and has long memory
- VIX
 - VIX from option prices on the SP500, as an approximate market estimate of the annualized volatility (std) of SP500 returns in the next month
 - Changes in volatility are strongly negatively correlated with SP500 returns
 - extremely skewed, could use log VIX (also skewed)
 - High persistence
 - Can change (especially increase) rapidly
 - has little predictive power for future returns
 - horizon ≈ 21
 - standardized returns: simple return/VIX, could assume normal or student-T
 - getting the distribution right can be more important than getting volatility right
 - VIX is a fairly good measure of volatility one month ahead
 - does not predict volatility at long horizons
 - low volatility breeds high volatility
 - we take a bigger position in SP500 when VIX is low, and then reduce the position when VIX increases
- two-stage analysis
- model volatility
 - realized volatility
 - $\sigma_t^2 = \frac{1}{h} \sum_{i=1}^h r_{t-i}^2$
 - $h=21$ for one month
 - Exponential smoothing (exponential moving averages)
 - $\sigma_t^2 = \frac{1}{\sum_{i=0}^{\infty} \delta^i} \sum_{i=1}^h \delta^i r_{t-i}^2$
 - $\sigma_t^2 = \delta \sigma_{t-1}^2 + (1 - \delta) r_{t-1}^2$
 - Two-stage analysis with the Risk Metrics approach
 - compute all σ^2
 - compute standardized return $z = (p_{t+1} - p_t) / \sigma_{t+1}$
 - forecast z
 - get forecast σ^2
 - estimate is biased: can use a MAD estimate to achieve a better fit
- Volatility clustering always generates thicker tails
- GARCH(1,1) model
- Threshold GARCH(1,1)

- that negative returns increase the variance more than positive returns
- negative return use big beta, positive return use small beta
- This gives $\text{corr}(\sigma_t+1, r_t) < 0$ and generates skewness in cumulative returns

9. Kelly criterion

- suggests to maximize the expected value of next period log wealth
- $f^*=2p-1$ (detail ★ see paper notes)
- maximize the median terminal wealth (while $f=1$ to maximize expected return)
- almost certain of increasing your wealth if you have an edge ($p>0.5$)
- $f=p/a-(1-p)/b$ (win b and loss a)
- Ruin impossible
- $f^* = \frac{\mu}{\sigma^2}$ (detail ★ see paper notes)
- our portfolio volatility is equal to the portfolio Sharpe ratio
- considering r_f , f^* the same
- half-Kelly
- mean-variance solution near Kelly's
 - The similarity to MV is only if the quadratic approximation holds. If ruin or extremely large losses are possible, the similarity is broken
 - The similarity is therefore most likely to hold when considering short horizons.
 - The Kelly criterion elucidates (point) the role of under-betting and over-betting, and the importance of avoiding ruin. This is typically lost in MV.
 - Most people and organizations cannot say what their λ is
- Much easier to specify the maximum drawdown (loss) they can tolerate.

10. multivariate modeling

- Multivariate normal
 - marginal distributions are also Gaussian
- Multivariate t
- Mixtures of multivariate normals and of multivariate t
- **Local correlations increasing in some periods (good) generates thicker(fatter) tails in portfolio returns**
- **If local correlations increase in bad periods, modest skew**
- A multivariate normal or even a multivariate t would miss this effect and produce too thin portfolio tails
- Factor modeling *
 - $\text{vcv}(R_t) = B[E(f_t f_t^T)]B^T + \sigma^2 I$
 - has many less parameters
 - factors may be explanatory variables like market returns, oil prices, GDP growth, or they may be statistical factors
- Risk-Metrics approach *
 - $V_t = \delta V_{t-1} + (1 - \delta) R_{t-1} R_{t-1}^T$

11. mean-variance optimization

- $\max(w) \quad [E(R_p) - \lambda \sigma^2(R_p)]$
- the w is the same as the approximate Kelly solution for f if $\lambda = 0.5$
- Both Kelly and MV are very sensitive to inputs
- (Kelly and MV) easy task due to
 - few assets
 - low or modest cross-correlations
 - knowledge of average returns (which is not available in real time)
- Problems with MV in practice
 - quasi(准)-arbitrage (like arbitrage, similar with arbitrage)
 - MV will try to use assets that arbitrage

- problems more severe with more assets
- errors in estimating $m_{t|t+1}$ are particularly detrimental to the performance of MV
- MV will work best with few assets (asset classes), modest(少量) correlations, and good estimates $m_{t|t+1}$ (knowledge of average returns)
- recent reactions to practical problems with MV
 - Equal weight portfolio (1/n)
 - Minimum-variance portfolio
 - Usually applied within an asset class, typically stocks
 - against the idea that risk should be rewarded
 - examples: high volatility != high return
 - Risk-parity portfolio
 - give the same volatility-weight to each asset class
 - target a desired level of portfolio volatility
 - Sorted portfolios
 - Other approaches
- Some relevant R packages *:
 1. fPortfolio. Large package covering many types of portfolio optimizations
 2. FRAPO. Draw-down constraints (and more)
 3. PerformanceAnalytics. Large package for measuring various portfolio statistics (VaR, ES, drawdowns)

12. Bayesian

- The Bayesian approach
 - Randomness is a state of knowledge, not a state of nature.
- base rates
 - we tend to over-weight the evidence $p(B|A)$ and neglect the base rate $p(A)$
 - The frequentist approach also neglects the base rate
 - neglecting the base rate can lead to large errors
 - $p(B|A)$ does not contain a lot information
 - at low signal-to-noise situations
 - In finance, low signal-to-noise is the rule
- Example: updating estimates of managers' skill (cannot understand)
- Bayesian inference(推导) on model parameters
- Forecasting *
- posterior distribution $p(\theta|y)$ (后验概率分布)
- Two types of priors (先验)
 - reflecting knowledge coming from experts' judgment, similar situations, or model's simulation
 - reflecting the idea that the final model will be in some form "simpler", in the sense that some elements of θ will be zero or small.
- Three types of inference *
 - available analytically
 - by computer intensive methods known as MCMC (Markov Chain Monte Carlo)
 - approximation
- Model averaging
 - consider more than one model, Bayesian inference taking a weighted combination of all the models
 - don't pick one, but average the results of both.

read: Rational Decision-Making Under Uncertainty

- introduce the coin experiment
- maybe you know Kelly criterion
- Findings: How Well Did Our Players Play
 - having heard of Kelly did not seem to help two of them
 - some take small and constant bet
 - gambler increases the size of his wagers after losses

- How Much Should You Be Willing to Pay to Play
 - depend on the cap and time
- Similarities to Investing in the Stock Market
- Conclusion

read: Superforecasting

- nothing new under the sun
- Superforecasters are not perfect Bayesian updaters but they are better than most
- generate different perspectives
- The more a forecaster embraced probabilistic thinking, the more accurate she was